

Data Visualization and Cleaning Assignment

Created by: Venkata Reddy Konasani

statinfer.com (<https://statinfer.com/>).

The end goal is to build a credit risk model that predicts the risky customers. In this assignment we are going to discuss data visualizations and cleaning process. Read the column names and descriptions before you start this assignment. Download the data from this below link

<https://drive.google.com/drive/folders/1ISpTlfUik4enxgDUkqPzNrDBevN9vuQN?usp=sharing>
(<https://drive.google.com/drive/folders/1ISpTlfUik4enxgDUkqPzNrDBevN9vuQN?usp=sharing>)

Note

Some questions may not be relevant for some variables. You can ignore questions that are not applicable.

1.Data Importing and basic Details

Import loans dataset. Print some sample rows and basic information. Read the data dictionary and understand the variables before starting this assignment

In [1]:

```
import pandas as pd
loans=pd.read_csv("/content/drive/My Drive/Training/ML_Full_Semester/Assignments/Data_v
iz_Cleaning_Give_me_Credit/Datasets/Loans/cs-training.csv")
```

In [2]:

```
print(loans.shape)
print(loans.columns)
```

```
(150000, 12)
Index(['Sr_No', 'SeriousDlqin2yrs', 'RevolvingUtilizationOfUnsecuredLine
s',
      'age', 'NumberOfTime30-59DaysPastDueNotWorse', 'DebtRatio',
      'MonthlyIncome', 'NumberOfOpenCreditLinesAndLoans',
      'NumberOfTimes90DaysLate', 'NumberRealEstateLoansOrLines',
      'NumberOfTime60-89DaysPastDueNotWorse', 'NumberOfDependents'],
      dtype='object')
```

In [3]:

```
loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150000 entries, 0 to 149999
Data columns (total 12 columns):
 #   Column                                                                 Non-Null Count  Dtype  
---  -
 0   Sr_No                                                                150000 non-null  int64  
 1   SeriousDlqin2yrs                                                    150000 non-null  int64  
 2   RevolvingUtilizationOfUnsecuredLines  150000 non-null  float64
 3   age                                                                  150000 non-null  int64  
 4   NumberOfTime30-59DaysPastDueNotWorse  150000 non-null  int64  
 5   DebtRatio                                                            150000 non-null  float64
 6   MonthlyIncome                                                        120269 non-null  float64
 7   NumberOfOpenCreditLinesAndLoans    150000 non-null  int64  
 8   NumberOfTimes90DaysLate          150000 non-null  int64  
 9   NumberRealEstateLoansOrLines       150000 non-null  int64  
10   NumberOfTime60-89DaysPastDueNotWorse  150000 non-null  int64  
11   NumberOfDependents              146076 non-null  float64
dtypes: float64(4), int64(8)
memory usage: 13.7 MB
```

In [4]:

```
loans.sample(5).T
```

Out[4]:

	149458	69106	20001	13492
Sr_No	149459.000000	69107.000000	20002.000000	134924.000000
SeriousDlqin2yrs	0.000000	0.000000	1.000000	0.000000
RevolvingUtilizationOfUnsecuredLines	0.365909	0.017334	0.885052	0.595700
age	61.000000	66.000000	70.000000	43.000000
NumberOfTime30-59DaysPastDueNotWorse	0.000000	0.000000	0.000000	0.000000
DebtRatio	0.461703	0.010278	0.836776	0.544400
MonthlyIncome	3733.000000	6615.000000	3436.000000	9900.000000
NumberOfOpenCreditLinesAndLoans	14.000000	8.000000	17.000000	10.000000
NumberOfTimes90DaysLate	0.000000	0.000000	0.000000	0.000000
NumberRealEstateLoansOrLines	0.000000	1.000000	1.000000	4.000000
NumberOfTime60-89DaysPastDueNotWorse	0.000000	0.000000	0.000000	0.000000
NumberOfDependents	0.000000	0.000000	0.000000	0.000000

Variable1 = "SeriousDlqin2yrs "

Is this a target column or a predictor column ?

In [5]:

```
#Target
```

What type of column is this? Numerical continuous, Numerical Discrete, DateTime, Geo

In [6]:

```
#Numerical discrete column  
loans["SeriousDlqin2yrs"].dtypes
```

Out[6]:

```
dtype('int64')
```

Perform Univariate Analysis. If it is continuous then histograms and box plots, if it is discrete or categorical then bar charts and frequency tables.

In [7]:

```
loans["SeriousDlqin2yrs"].value_counts()
```

Out[7]:

```
0    139974  
1     10026  
Name: SeriousDlqin2yrs, dtype: int64
```

In [8]:

```
loans["SeriousDlqin2yrs"].value_counts()
```

Out[8]:

```
0    139974  
1     10026  
Name: SeriousDlqin2yrs, dtype: int64
```


In [9]:

```
import seaborn as sns
sns.countplot(x="SeriousDlqin2yrs", data=loans)
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
  import pandas.util.testing as tm
```

Out[9]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bd8cda90>



List down the percentage of missing and percentage of lowside and high side outliers

In [10]:

```
loans["SeriousDlqin2yrs"].isnull().sum()
```

Out[10]:

0

Validate this variable, Perform some data checks. If you find any issues then write down the issues.

In [11]:

```
#No issues
```

If you find any issues then clean the variable, by applying appropriate imputation or any other substitution technique.

In [12]:

```
#NA
```

Create a cross tab or pivot table with the target variable and check the power of prediction. See if there are any apparent patterns in this variable

In [13]:

```
#NA
```

If you have any additional findings, that are missing above, highlight them here

In [14]:

```
#NA
```

Variable2 = "RevolvingUtilizationOfUnsecuredLines"

Is this a target column or a predictor column ?

In [15]:

```
#Predictor
```

What type of column is this? Numerical continuous, Numerical Discrete, DateTime, Geo

In [16]:

```
# Numerical continuous  
loans["RevolvingUtilizationOfUnsecuredLines"].dtypes
```

Out[16]:

```
dtype('float64')
```

Perform Univariate Analysis. If it is continuous then histograms and box plots, if it is discrete or categorical then bar charts and frequency tables.

In [17]:

```
loans["RevolvingUtilizationOfUnsecuredLines"].describe()
```

Out[17]:

```
count    150000.000000
mean         6.048438
std       249.755371
min         0.000000
25%        0.029867
50%        0.154181
75%        0.559046
max       50708.000000
Name: RevolvingUtilizationOfUnsecuredLines, dtype: float64
```

In [18]:

```
loans["RevolvingUtilizationOfUnsecuredLines"].quantile([0,0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,0.91,0.92,0.93,0.94,0.95,0.96,0.97,0.98,0.99,1])
```

Out[18]:


```
0.00    0.000000
0.01    0.000000
0.02    0.000000
0.03    0.000000
0.04    0.000000
0.05    0.000000
0.06    0.000000
0.07    0.000000
0.08    0.000708
0.09    0.001733
0.10    0.002969
0.20    0.019222
0.30    0.043461
0.40    0.083181
0.50    0.154181
0.60    0.271493
0.70    0.445136
0.80    0.698857
0.90    0.981278
0.91    1.000000
0.92    1.000000
0.93    1.000000
0.94    1.000000
0.95    1.000000
0.96    1.000000
0.97    1.000000
0.98    1.006199
0.99    1.092956
1.00    50708.000000
Name: RevolvingUtilizationOfUnsecuredLines, dtype: float64
```

In [19]:

```
sns.boxplot(y=loans["RevolvingUtilizationOfUnsecuredLines"])
```

Out[19]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bd2fb7b8>




In [20]:

```
sns.distplot(loans["RevolvingUtilizationOfUnsecuredLines"])
```

Out[20]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bd2fb390>



List down the percentage of missing and percentage of lowside and high side outliers

In [21]:

```
loans["RevolvingUtilizationOfUnsecuredLines"].isnull().sum()
```

Out[21]:

0

In [22]:

```
#3% high side outliers
```

Validate this variable, Perform some data checks. If you find any issues then write down the issues.

In [23]:

```
# 3% issues - No missing values
```

If you find any issues then clean the variable, by applying appropriate imputation or any other substitution technique.

In [24]:

```
loans["util_new"]=loans["RevolvingUtilizationOfUnsecuredLines"]  
loans["util_new"][loans["util_new"]>1]=loans["util_new"].median()
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

In [25]:

```
loans["util_new"].describe()
```

Out[25]:

```
count    150000.000000  
mean         0.300469  
std         0.334855  
min         0.000000  
25%         0.029867  
50%         0.154178  
75%         0.506929  
max         1.000000  
Name: util_new, dtype: float64
```

Create a cross tab or pivot table with the target variable and check the power of prediction. See if there are any apparent patterns in this variable

In [26]:

```
util_pivot=pd.pivot_table(data=loans, values='util_new', columns="SeriousDlqin2yrs", aggfunc='mean')  
print(util_pivot)
```


```
SeriousDlqin2yrs      0      1  
util_new      0.280592  0.577982
```


In [27]:

```
sns.boxplot(x=loans["SeriousDlqin2yrs"], y=loans["util_new"])
```

Out[27]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bce5aba8>




In [28]:

```
sns.distplot(loans[loans["SeriousDlqin2yrs"]==1]["util_new"], color="red")  
sns.distplot(loans[loans["SeriousDlqin2yrs"]==0]["util_new"], color="green")
```

Out[28]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bcd299e8>




In [29]:

```
sns.violinplot(x=loans["SeriousDlqin2yrs"], y=loans["util_new"])
```

Out[29]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bcbc2f98>



If you have any additional findings, that are missing above, highlight them here

In [30]:

```
# Util is high when serious delinquency=1
```

Variable3= "age"

Is this a target column or a predictor column ?

In [31]:

```
#Predictor
```

What type of column is this? Numerical continuous, Numerical Discrete, DateTime, Geo

In [32]:

```
#Numeric Discrete
```

Perform Univariate Analysis. If it is continuous then histograms and box plots, if it is discrete or categorical then bar charts and frequency tables.

In [33]:

```
loans["age"].value_counts()
```

Out[33]:

```
49    3837
48    3806
50    3753
63    3719
47    3719
...
101     3
109     2
107     1
105     1
0       1
Name: age, Length: 86, dtype: int64
```

In [34]:

```
loans["age"].describe()
```

Out[34]:


```
count    150000.000000
mean       52.295207
std       14.771866
min        0.000000
25%       41.000000
50%       52.000000
75%       63.000000
max      109.000000
Name: age, dtype: float64
```

In [35]:

```
sns.boxplot(y=loans["age"])
```

Out[35]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bccb1320>




In [36]:

```
sns.distplot(loans["age"])
```

Out[36]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17ba300828>



List down the percentage of missing and percentage of lowside and high side outliers

In [37]:

```
loans["age"].quantile([0,0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,0.91,0.92,0.93,0.94,0.95,0.96,0.97,0.98,0.99,1])
```

Out[37]:

```
0.00    0.0
0.01    24.0
0.02    25.0
0.03    27.0
0.04    28.0
0.05    29.0
0.06    30.0
0.07    30.0
0.08    31.0
0.09    32.0
0.10    33.0
0.20    39.0
0.30    44.0
0.40    48.0
0.50    52.0
0.60    56.0
0.70    61.0
0.80    65.0
0.90    72.0
0.91    73.0
0.92    74.0
0.93    75.0
0.94    76.0
0.95    78.0
0.96    79.0
0.97    81.0
0.98    84.0
0.99    87.0
1.00   109.0
Name: age, dtype: float64
```

Validate this variable, Perform some data checks. If you find any issues then write down the issues.

In [38]:

```
#1 % high side outliers
```

If you find any issues then clean the variable, by applying appropriate imputation or any other substitutaion technique.

In [39]:

```
loans["age_new"]=loans["age"]
loans["age_new"][loans["age_new"]>90]=loans["age"].median()
sns.boxplot(y=loans["age_new"])
```


/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[39]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17bcd92668>



Create a cross tab or pivot table with the target variable and check the power of prediction. See if there are any apparent patterns in this variable

In [40]:

```
age_pivot=pd.pivot_table(data=loans, values='age_new', columns="SeriousDlqin2yrs", aggfunc='mean')
print(age_pivot)
```


SeriousDlqin2yrs	0	1
age_new	52.610878	45.880311

In [41]:

```
sns.boxplot(x=loans["SeriousDlqin2yrs"], y=loans["age_new"])
```

Out[41]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17ba1d3390>




In [42]:

```
sns.distplot(loans[loans["SeriousDlqin2yrs"]==1]["age_new"], color="red")  
sns.distplot(loans[loans["SeriousDlqin2yrs"]==0]["age_new"], color="green")
```

Out[42]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17ba14e080>




In [43]:

```
sns.violinplot(x=loans["SeriousDlqin2yrs"], y=loans["age_new"])
```

Out[43]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b9faecc0>



If you have any additional findings, that are missing above, highlight them here

In [44]:

```
#NA
```

Variable4 = "NumberOfTime30-59DaysPastDueNotWorse"

Is this a target column or a predictor column ?

In [45]:

```
#predictor
```

What type of column is this? Numerical continuous, Numerical Discrete, DateTime, Geo

In [46]:

```
#Numerical discrete
```

Perform Univariate Analysis. If it is continuous then histograms and box plots, if it is discrete or categorical then bar charts and frequency tables.

In [47]:

```
loans["NumberOfTime30-59DaysPastDueNotWorse"].value_counts()
```

Out[47]:


```
0      126018
1       16033
2        4598
3        1754
4         747
5         342
98        264
6         140
7          54
8          25
9          12
96           5
10           4
12           2
13           1
11           1
Name: NumberOfTime30-59DaysPastDueNotWorse, dtype: int64
```

In [48]:

```
sns.countplot(x="NumberOfTime30-59DaysPastDueNotWorse", data=loans)
```

Out[48]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b9f80860>



List down the percentage of missing and percentage of lowside and high side outliers

In [49]:

```
loans["NumberOfTime30-59DaysPastDueNotWorse"].value_counts()/len(loans)
```

Out[49]:

```
0      0.840120
1      0.106887
2      0.030653
3      0.011693
4      0.004980
5      0.002280
98      0.001760
6      0.000933
7      0.000360
8      0.000167
9      0.000080
96      0.000033
10      0.000027
12      0.000013
13      0.000007
11      0.000007
Name: NumberOfTime30-59DaysPastDueNotWorse, dtype: float64
```

Validate this variable, Perform some data checks. If you find any issues then write down the issues.

In [50]:

```
#Less than 1% have issues
```

If you find any issues then clean the variable, by applying appropriate imputation or any other substitutaion technique.

In [51]:

```
cross_tab_30dpd_target=pd.crosstab(loans[ 'NumberOfTime30-59DaysPastDueNotWorse' ],loans[
'SeriousDlqin2yrs'])
cross_tab_30dpd_target
```

Out[51]:

SeriousDlqin2yrs		0	1
NumberOfTime30-59DaysPastDueNotWorse			
	0	120977	5041
	1	13624	2409
	2	3379	1219
	3	1136	618
	4	429	318
	5	188	154
	6	66	74
	7	26	28
	8	17	8
	9	8	4
	10	1	3
	11	0	1
	12	1	1
	13	0	1
	96	1	4
	98	121	143

In [52]:

```
cross_tab_30dpd_target_percent=cross_tab_30dpd_target.astype(float).div(cross_tab_30dpd_target.sum(axis=1), axis=0)
round(cross_tab_30dpd_target_percent,2)
```

Out[52]:

	SeriousDlqin2yrs	0	1
NumberOfTime30-59DaysPastDueNotWorse			
	0	0.96	0.04
	1	0.85	0.15
	2	0.73	0.27
	3	0.65	0.35
	4	0.57	0.43
	5	0.55	0.45
	6	0.47	0.53
	7	0.48	0.52
	8	0.68	0.32
	9	0.67	0.33
	10	0.25	0.75
	11	0.00	1.00
	12	0.50	0.50
	13	0.00	1.00
	96	0.20	0.80
	98	0.46	0.54

In [53]:

```
loans['num_30_59_dpd_new']=loans['NumberOfTime30-59DaysPastDueNotWorse']
loans['num_30_59_dpd_new'][loans['num_30_59_dpd_new']>12]=6
loans['num_30_59_dpd_new']

loans['num_30_59_dpd_new'].value_counts(sort=False)
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[53]:

```
0    126018
1     16033
2     4598
3     1754
4      747
5     342
6     410
7      54
8      25
9      12
10      4
11      1
12      2
Name: num_30_59_dpd_new, dtype: int64
```


Create a cross tab or pivot table with the target variable and check the power of prediction. See if there are any apparent patterns in this variable

In [54]:

```
sns.countplot(loans[loans["SeriousDlqin2yrs"]==0]["num_30_59_dpd_new"], color="red")
sns.countplot(loans[loans["SeriousDlqin2yrs"]==1]["num_30_59_dpd_new"], color="green")
```

Out[54]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b9c46710>

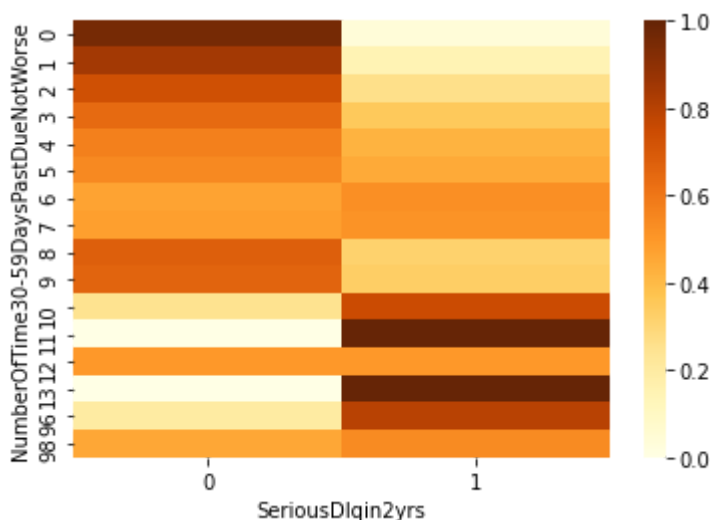


In [55]:

```
cross_tab_30dpd_target_percent=cross_tab_30dpd_target.astype(float).div(cross_tab_30dpd_target.sum(axis=1), axis=0)
round(cross_tab_30dpd_target_percent,2)
sns.heatmap(cross_tab_30dpd_target_percent, cmap="YlOrBr")
```

Out[55]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b9ee49e8>



If you have any additional findings, that are missing above, highlight them here

In [56]:

```
# As the 30DPD incrases the delinquencies are increasing
```

Variable5 = "DebtRatio"

Is this a target column or a predictor column ?

In [57]:

```
#Predictor
```

What type of column is this? Numerical continuous, Numerical Discrete, DateTime, Geo

In [58]:

```
# Numerical continuous
loans["DebtRatio"].dtypes
```

Out[58]:

```
dtype('float64')
```

Perform Univariate Analysis. If it is continuous then histograms and box plots, if it is discrete or categorical then bar charts and frequency tables.

In [59]:

```
loans["DebtRatio"].describe()
```

Out[59]:

```
count    150000.000000
mean       353.005076
std       2037.818523
min         0.000000
25%        0.175074
50%        0.366508
75%        0.868254
max       329664.000000
Name: DebtRatio, dtype: float64
```

In [60]:

```
loans["DebtRatio"].quantile([0,0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.75,0.76,0.77, 0.78,0.8,0.81,0.82, 0.83,0.86,0.90,0.93,0.94,0.95,0.96,0.97,0.98,0.99,1])
```

Out[60]:

0.00	0.000000
0.01	0.000000
0.02	0.000000
0.03	0.000480
0.04	0.002278
0.05	0.004329
0.06	0.006799
0.07	0.009997
0.08	0.014780
0.09	0.021249
0.10	0.030874
0.20	0.133773
0.30	0.213697
0.40	0.287460
0.50	0.366508
0.60	0.467506
0.70	0.649189
0.75	0.868254
0.76	0.951184
0.77	1.058832
0.78	1.275069
0.80	4.000000
0.81	14.000000
0.82	31.000000
0.83	61.000000
0.86	453.000000
0.90	1267.000000
0.93	1917.070000
0.94	2172.060000
0.95	2449.000000
0.96	2791.000000
0.97	3225.000000
0.98	3839.000000
0.99	4979.040000
1.00	329664.000000


Name: DebtRatio, dtype: float64

In [61]:

```
sns.boxplot(y=loans["DebtRatio"])
```

Out[61]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17ba1f3588>




In [62]:

```
sns.distplot(loans["DebtRatio"])
```

Out[62]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17ba1378d0>



List down the percentage of missing and percentage of lowside and high side outliers

In [63]:

```
loans["DebtRatio"].isnull().sum()
```

Out[63]:

0

In [64]:

```
#24% outliers  
#No missing values
```

Validate this variable, Perform some data checks. If you find any issues then write down the issues.

In [65]:

```
#24% outliers  
#No missing values
```

If you find any issues then clean the variable, by applying appropriate imputation or any other substitution technique.

In [66]:

```
loans["DebtRatio_new_ind"]=(loans["DebtRatio"]>1)*1  
loans["DebtRatio_new_ind"].value_counts()
```

Out[66]:

```
0    114863  
1     35137  
Name: DebtRatio_new_ind, dtype: int64
```

In [67]:

```
loans["DebtRatio_new"]=loans["DebtRatio"]  
loans["DebtRatio_new"][loans["DebtRatio_new"]>1]=loans["DebtRatio"].median()
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

In [68]:

```
loans["DebtRatio_new"].describe()
```

Out[68]:


```
count    150000.000000
mean       0.317893
std        0.199835
min         0.000000
25%        0.175074
50%        0.366506
75%        0.380021
max         1.000000
Name: DebtRatio_new, dtype: float64
```

In [69]:

```
sns.boxplot(y=loans["DebtRatio_new"])
```

Out[69]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b8979160>




In [70]:

```
sns.distplot(loans["DebtRatio_new"])
```

Out[70]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b8958a90>



Create a cross tab or pivot table with the target variable and check the power of prediction. See if there are any apparent patterns in this variable

In [71]:

```
Debt_ratio_pivot=pd.pivot_table(data=loans, values='DebtRatio_new', columns="SeriousDlqin2yrs", aggfunc='mean')  
print(Debt_ratio_pivot)
```


```
SeriousDlqin2yrs      0      1  
DebtRatio_new      0.31543  0.352284
```

In [72]:

```
sns.boxplot(x=loans["SeriousDlqin2yrs"], y=loans["DebtRatio_new"])
```

Out[72]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b884def0>

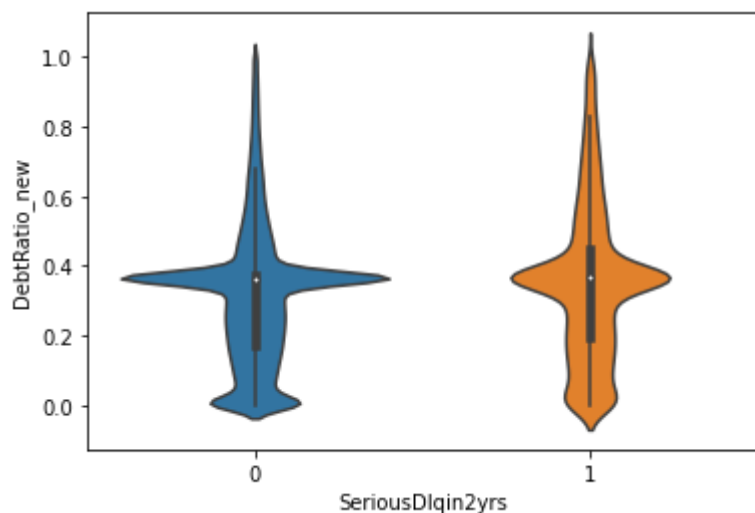


In [73]:

```
sns.violinplot(x=loans["SeriousDlqin2yrs"], y=loans["DebtRatio_new"])
```

Out[73]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b87bfa58>



If you have any additional findings, that are missing above, highlight them here

In [74]:

```
#NA
```

Variable6 = "MonthlyIncome"

Is this a target column or a predictor column ?

In [75]:

```
#Predictor
```

What type of column is this? Numerical continuous, Numerical Discrete, DateTime, Geo

In [76]:

```
# Numerical continuous  
loans["MonthlyIncome"].dtypes
```

Out[76]:

```
dtype('float64')
```

Perform Univariate Analysis. If it is continuous then histograms and box plots, if it is discrete or categorical then bar charts and frequency tables.

In [77]:

```
loans["MonthlyIncome"].describe().round()
```

Out[77]:

```
count      120269.0  
mean        6670.0  
std         14385.0  
min           0.0  
25%         3400.0  
50%         5400.0  
75%         8249.0  
max        3008750.0  
Name: MonthlyIncome, dtype: float64
```

In [78]:

```
loans["MonthlyIncome"].quantile([0,0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09,0.1,0.2,0.3,0.4,0.5,0.6,0.7, 0.83,0.86,0.90,0.93,0.94,0.95,0.96,0.97,0.98,0.99,1])
```

Out[78]:


```
0.00      0.00
0.01      0.00
0.02     250.00
0.03     800.00
0.04    1012.16
0.05    1300.00
0.06    1500.00
0.07    1640.00
0.08    1800.00
0.09    1981.00
0.10    2005.00
0.20    3000.00
0.30    3800.00
0.40    4544.20
0.50    5400.00
0.60    6300.00
0.70    7500.00
0.83    9800.00
0.86   10417.00
0.90   11666.00
0.93   13000.00
0.94   13716.00
0.95   14587.60
0.96   15636.28
0.97   17000.00
0.98   19600.00
0.99   25000.00
1.00  3008750.00
Name: MonthlyIncome, dtype: float64
```

In [79]:

```
sns.boxplot(y=loans["MonthlyIncome"])
```

Out[79]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b87a1048>




In [80]:

```
sns.distplot(loans["MonthlyIncome"])
```

Out[80]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b8710668>



List down the percentage of missing and percentage of lowside and high side outliers

In [81]:

```
loans["MonthlyIncome"].isnull().sum()
```

Out[81]:

29731

In [82]:

```
loans["MonthlyIncome"].isnull().sum()/len(loans)
```

Out[82]:

0.19820666666666667

Validate this variable, Perform some data checks. If you find any issues then write down the issues.

In [83]:

```
#2% igh side outliers
```

In [84]:

```
#19.8% missing values
```

If you find any issues then clean the variable, by applying appropriate imputation or any other substitutaion technique.

In [85]:

```
loans['MonthlyIncome_ind']=1
loans['MonthlyIncome_ind'][loans['MonthlyIncome'].isnull()]=0
loans['MonthlyIncome_ind'].value_counts(sort=False)
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[85]:

```
0      29731
1     120269
Name: MonthlyIncome_ind, dtype: int64
```

In [86]:

```
loans['MonthlyIncome_new']=loans['MonthlyIncome']
loans['MonthlyIncome_new'][loans['MonthlyIncome'].isnull()]=loans['MonthlyIncome'].median()
round(loans['MonthlyIncome_new'].describe())
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[86]:

```
count      150000.0
mean         6418.0
std         12890.0
min           0.0
25%          3903.0
50%          5400.0
75%          7400.0
max        3008750.0
Name: MonthlyIncome_new, dtype: float64
```


In [87]:

```
loans['MonthlyIncome_new'][loans['MonthlyIncome']>20000]=loans['MonthlyIncome'].median()  
round(loans['MonthlyIncome_new'].describe())
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""Entry point for launching an IPython kernel.

Out[87]:

```
count    150000.0  
mean       5902.0  
std       3298.0  
min         0.0  
25%       3903.0  
50%       5400.0  
75%       7100.0  
max      20000.0  
Name: MonthlyIncome_new, dtype: float64
```

Create a cross tab or pivot table with the target variable and check the power of prediction. See if there are any apparent patterns in this variable

In [88]:

```
MonthlyIncome_pivot=pd.pivot_table(data=loans, values='MonthlyIncome_new', columns="SeriousDlqin2yrs", aggfunc='mean')  
print(MonthlyIncome_pivot)
```


```
SeriousDlqin2yrs      0      1  
MonthlyIncome_new  5949.503372  5237.481049
```

In [89]:

```
sns.boxplot(y=loans["MonthlyIncome_new"])
```

Out[89]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b8608e80>




In [90]:

```
sns.violinplot(x=loans["SeriousDlqin2yrs"], y=loans["MonthlyIncome_new"])
```

Out[90]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b8571f60>



If you have any additional findings, that are missing above, highlight them here

In [91]:

```
#Monthly income seems to have no impact on the overall delinquency
```

Variable7 = "NumberOfOpenCreditLinesAndLoans"

Is this a target column or a predictor column ?

In [92]:

```
#Predictor
```

What type of column is this? Numerical continuous, Numerical Discrete, DateTime, Geo

In [93]:

```
#Numerical discrete  
loans["NumberOfOpenCreditLinesAndLoans"].dtypes
```

Out[93]:

```
dtype('int64')
```

Perform Univariate Analysis. If it is continuous then histograms and box plots, if it is discrete or categorical then bar charts and frequency tables.

In [94]:

```
loans["NumberOfOpenCreditLinesAndLoans"].value_counts(sort=False)
```

Out[94]:

0	1888
1	4438
2	6666
3	9058
4	11609
5	12931
6	13614
7	13245
8	12562
9	11355
10	9624
11	8321
12	7005
13	5667
14	4546
15	3645
16	3000
17	2370
18	1874
19	1433
20	1169
21	864
22	685
23	533
24	422
25	337
26	239
27	194
28	150
29	114
30	88
31	74
32	52
33	47
34	35
35	27
36	18
37	7
38	13
39	9
40	10
41	4
42	8
43	8
44	2
45	8
46	3
47	2
48	6
49	4
50	2
51	2
52	3
53	1
54	4
56	2
57	2
58	1


Name: NumberOfOpenCreditLinesAndLoans, dtype: int64

In [95]:

```
sns.countplot(x="NumberOfOpenCreditLinesAndLoans", data=loans)
```

Out[95]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b84fe208>




In [96]:

```
sns.countplot(y="NumberOfOpenCreditLinesAndLoans", data=loans)
```

Out[96]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17ba1ebcf8>



In [97]:

```
loans["NumberOfOpenCreditLinesAndLoans"].value_counts()/len(loans)
```

Out[97]:

6	0.090760
7	0.088300
5	0.086207
8	0.083747
4	0.077393
9	0.075700
10	0.064160
3	0.060387
11	0.055473
12	0.046700
2	0.044440
13	0.037780
14	0.030307
1	0.029587
15	0.024300
16	0.020000
17	0.015800
0	0.012587
18	0.012493
19	0.009553
20	0.007793
21	0.005760
22	0.004567
23	0.003553
24	0.002813
25	0.002247
26	0.001593
27	0.001293
28	0.001000
29	0.000760
30	0.000587
31	0.000493
32	0.000347
33	0.000313
34	0.000233
35	0.000180
36	0.000120
38	0.000087
40	0.000067
39	0.000060
45	0.000053
43	0.000053
42	0.000053
37	0.000047
48	0.000040
41	0.000027
54	0.000027
49	0.000027
46	0.000020
52	0.000020
51	0.000013
56	0.000013
57	0.000013
50	0.000013
47	0.000013
44	0.000013
53	0.000007
58	0.000007

Name: NumberOfOpenCreditLinesAndLoans, dtype: float64

In [98]:

```
loans["NumberOfOpenCreditLinesAndLoans"].describe()
```

Out[98]:


```
count    150000.000000
mean         8.452760
std         5.145951
min          0.000000
25%          5.000000
50%          8.000000
75%         11.000000
max         58.000000
Name: NumberOfOpenCreditLinesAndLoans, dtype: float64
```

In [99]:

```
sns.boxplot(y=loans["NumberOfOpenCreditLinesAndLoans"])
```

Out[99]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b8549e48>



List down the percentage of missing and percentage of lowside and high side outliers

In [100]:

```
loans["NumberOfOpenCreditLinesAndLoans"].isnull().sum()
```

Out[100]:

0

In [101]:

```
loans["NumberOfOpenCreditLinesAndLoans"].quantile([0,0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,0.91,0.92,0.93,0.94,0.95,0.96,0.97,0.98,0.99,1])
```

Out[101]:

0.00	0.0
0.01	0.0
0.02	1.0
0.03	1.0
0.04	1.0
0.05	2.0
0.06	2.0
0.07	2.0
0.08	2.0
0.09	3.0
0.10	3.0
0.20	4.0
0.30	5.0
0.40	6.0
0.50	8.0
0.60	9.0
0.70	10.0
0.80	12.0
0.90	15.0
0.91	16.0
0.92	16.0
0.93	17.0
0.94	17.0
0.95	18.0
0.96	19.0
0.97	20.0
0.98	22.0
0.99	24.0
1.00	58.0

Name: NumberOfOpenCreditLinesAndLoans, dtype: float64

Validate this variable, Perform some data checks. If you find any issues then write down the issues.

In [102]:

```
#3% High side outliers  
#No Missing values
```

If you find any issues then clean the variable, by applying appropriate imputation or any other substitutaion technique.

In [103]:

```
loans["Open_Credit_lines_new"]=loans["NumberOfOpenCreditLinesAndLoans"]
loans["Open_Credit_lines_new"][loans["Open_Credit_lines_new"]>20]=loans["NumberOfOpenCr
editLinesAndLoans"].median()
sns.boxplot(y=loans["Open_Credit_lines_new"])
```


/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[103]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b81d3a58>



Create a cross tab or pivot table with the target variable and check the power of prediction. See if there are any apparent patterns in this variable

In [104]:

```
credit_lines_pivot=pd.pivot_table(data=loans, values='Open_Credit_lines_new', columns=
"SeriousDlqin2yrs", aggfunc='mean')
print(credit_lines_pivot)
```


SeriousDlqin2yrs	0	1
Open_Credit_lines_new	8.05251	7.408039

In [105]:

```
sns.boxplot(x=loans["SeriousDlqin2yrs"], y=loans["Open_Credit_lines_new"])
```

Out[105]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b81399e8>

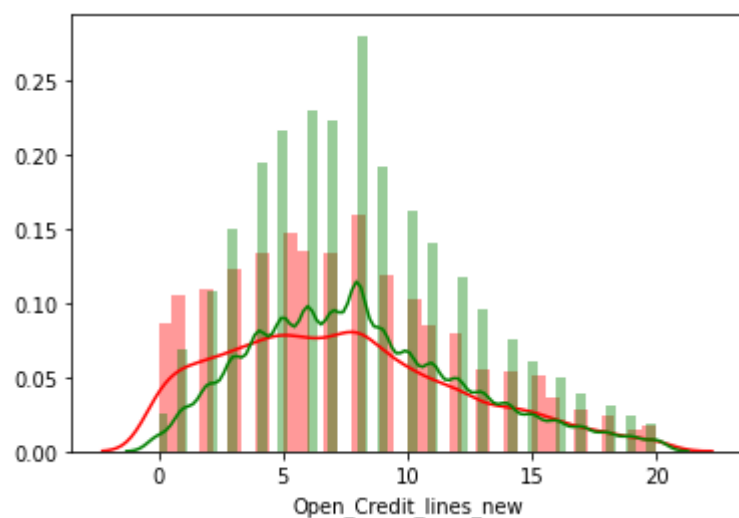


In [106]:

```
sns.distplot(loans[loans["SeriousDlqin2yrs"]==1]["Open_Credit_lines_new"], color="red")  
sns.distplot(loans[loans["SeriousDlqin2yrs"]==0]["Open_Credit_lines_new"], color="green")
```

Out[106]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b80ee780>




In [107]:

```
sns.violinplot(x=loans["SeriousDlqin2yrs"], y=loans["Open_Credit_lines_new"])
```

Out[107]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b7f8e208>



If you have any additional findings, that are missing above, highlight them here

In [108]:

```
#Looks like Open credit lines has no direct impact on the target
```

Variable8 = "NumberOfTimes90DaysLate"

This variable is similar to 30_59_DPD. Directly perform cleaning of this variable with very few lines of code. No need of visualization

In [109]:

```
cross_tab_90dpd_target=pd.crosstab(loans['NumberOfTimes90DaysLate'],loans['SeriousDlqin2yrs'])
cross_tab_90dpd_target
```

Out[109]:

SeriousDlqin2yrs		0	1
NumberOfTimes90DaysLate			
0	135108	6554	
1	3478	1765	
2	779	776	
3	282	385	
4	96	195	
5	48	83	
6	32	48	
7	7	31	
8	6	15	
9	5	14	
10	3	5	
11	2	3	
12	1	1	
13	2	2	
14	1	1	
15	2	0	
17	0	1	
96	1	4	
98	121	143	

In [110]:

```
cross_tab_90dpd_target_percent=cross_tab_90dpd_target.astype(float).div(cross_tab_90dpd_target.sum(axis=1), axis=0)
round(cross_tab_90dpd_target_percent,2)
```

Out[110]:

SeriousDlqin2yrs	0	1
NumberOfTimes90DaysLate		
0	0.95	0.05
1	0.66	0.34
2	0.50	0.50
3	0.42	0.58
4	0.33	0.67
5	0.37	0.63
6	0.40	0.60
7	0.18	0.82
8	0.29	0.71
9	0.26	0.74
10	0.38	0.62
11	0.40	0.60
12	0.50	0.50
13	0.50	0.50
14	0.50	0.50
15	1.00	0.00
17	0.00	1.00
96	0.20	0.80
98	0.46	0.54

In [111]:

```
loans['num_90_dpd_new']=loans['NumberOfTimes90DaysLate']
loans['num_90_dpd_new'][loans['num_90_dpd_new']>12]=3
loans['num_90_dpd_new']

loans['num_90_dpd_new'].value_counts(sort=False)
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[111]:

```
0    141662
1     5243
2     1555
3      945
4      291
5      131
6       80
7       38
8       21
9       19
10      8
11      5
12      2
Name: num_90_dpd_new, dtype: int64
```

Variable9 = "NumberRealEstateLoansOrLines"

This variable is similar to Number of Open Credit Lines and Loans. Directly perform cleaning of this variable with very few lines of code. No need of visualization

In [112]:

```
loans["NumberRealEstateLoansOrLines"].describe()
```

Out[112]:

```
count    150000.000000
mean         1.018240
std         1.129771
min         0.000000
25%         0.000000
50%         1.000000
75%         2.000000
max        54.000000
Name: NumberRealEstateLoansOrLines, dtype: float64
```

In [113]:

```
loans["NumberRealEstateLoansOrLines"].quantile([0,0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,0.91,0.92,0.93,0.94,0.95,0.96,0.97,0.98,0.99,1])
```

Out[113]:

0.00	0.0
0.01	0.0
0.02	0.0
0.03	0.0
0.04	0.0
0.05	0.0
0.06	0.0
0.07	0.0
0.08	0.0
0.09	0.0
0.10	0.0
0.20	0.0
0.30	0.0
0.40	1.0
0.50	1.0
0.60	1.0
0.70	1.0
0.80	2.0
0.90	2.0
0.91	2.0
0.92	2.0
0.93	2.0
0.94	3.0
0.95	3.0
0.96	3.0
0.97	3.0
0.98	4.0
0.99	4.0
1.00	54.0

Name: NumberRealEstateLoansOrLines, dtype: float64

In [114]:

```
loans["Real_estate_loans_new"]=loans["NumberRealEstateLoansOrLines"]
loans["Real_estate_loans_new"][loans["Real_estate_loans_new"]>4]=loans["NumberRealEstateLoansOrLines"].median()
sns.boxplot(y=loans["Real_estate_loans_new"])
```


/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[114]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b7ee7438>



Variable10 = "NumberOfTime60-89DaysPastDueNotWorse"

This variable is similar to 30_59_DPD and 90_DPD. Directly perform cleaning of this variable with very few lines of code. No need of visualization

In [115]:

```
cross_tab_60dpd_target=pd.crosstab(loans[ 'NumberOfTime60-89DaysPastDueNotWorse' ],loans[
'SeriousDlqin2yrs'])
cross_tab_60dpd_target
```

Out[115]:

SeriousDlqin2yrs		0	1
NumberOfTime60-89DaysPastDueNotWorse			
0		135140	7256
1		3954	1777
2		557	561
3		138	180
4		40	65
5		13	21
6		4	12
7		4	5
8		1	1
9		1	0
11		0	1
96		1	4
98		121	143

In [116]:

```
cross_tab_60dpd_target_percent=cross_tab_60dpd_target.astype(float).div(cross_tab_60dpd_target.sum(axis=1), axis=0)
round(cross_tab_60dpd_target_percent,2)
```

Out[116]:

	SeriousDlqin2yrs	0	1
NumberOfTime60-89DaysPastDueNotWorse			
0	0.95	0.05	
1	0.69	0.31	
2	0.50	0.50	
3	0.43	0.57	
4	0.38	0.62	
5	0.38	0.62	
6	0.25	0.75	
7	0.44	0.56	
8	0.50	0.50	
9	1.00	0.00	
11	0.00	1.00	
96	0.20	0.80	
98	0.46	0.54	

In [117]:

```
loans['num_60_dpd_new']=loans['NumberOfTime60-89DaysPastDueNotWorse']
loans['num_60_dpd_new'][loans['num_60_dpd_new']>12]=3
loans['num_60_dpd_new']

loans['num_60_dpd_new'].value_counts(sort=False)
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[117]:

```
0      142396
1       5731
2       1118
3        587
4        105
5         34
6         16
7          9
8          2
9          1
10         1
Name: num_60_dpd_new, dtype: int64
```

Variable11 = "NumberOfDependents"

This variable is similar to Number of Open Credit Lines and Real estate loans. Directly perform cleaning of this variable with very few lines of code. No need of visualization

In [118]:

```
loans["NumberOfDependents"].describe()
```

Out[118]:

```
count      146076.000000
mean         0.757222
std         1.115086
min          0.000000
25%          0.000000
50%          0.000000
75%          1.000000
max         20.000000
Name: NumberOfDependents, dtype: float64
```

In [119]:

```
loans["NumberOfDependents"].quantile([0,0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,0.91,0.92,0.93,0.94,0.95,0.96,0.97,0.98,0.99,1])
```

Out[119]:

0.00	0.0
0.01	0.0
0.02	0.0
0.03	0.0
0.04	0.0
0.05	0.0
0.06	0.0
0.07	0.0
0.08	0.0
0.09	0.0
0.10	0.0
0.20	0.0
0.30	0.0
0.40	0.0
0.50	0.0
0.60	1.0
0.70	1.0
0.80	2.0
0.90	2.0
0.91	3.0
0.92	3.0
0.93	3.0
0.94	3.0
0.95	3.0
0.96	3.0
0.97	3.0
0.98	4.0
0.99	4.0
1.00	20.0

Name: NumberOfDependents, dtype: float64

In [120]:

```
loans["NumberOfDependents_new"] = loans["NumberOfDependents"]
loans["NumberOfDependents_new"][loans["NumberOfDependents_new"] > 8] = loans["NumberOfDependents"].median()
sns.boxplot(y=loans["NumberOfDependents_new"])
```


/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[120]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f17b7ed7d68>



Write down your final findings and treatment in two lines.

SeriousDlqin2yrs ==> Findings and Treatment

In [120]:

RevolvingUtilizationOfUnsecuredLines ==> Findings and Treatment

In [120]:

age ==> Findings and Treatment

In [120]:

NumberOfTime30-59DaysPastDueNotWorse ==> Findings and Treatmet

In [120]:

Debt Ratio ==> Findings and Treatmet

In [120]:

Monthly Income ==> Findings and Treatmet

In [120]:

NumberOfOpenCreditLines and Loans ==> Findings and Treatmet

In [120]:

NumberOfTimes90DaysLate ==> Findings and Treatmet

In [120]:

NumberRealEstateLoansOrLines ==> Findings and Treatmet

In [120]:

NumberOfTime60-89DaysPastDueNotWorse ==> Findings and Treatmet

In [120]:

NumberOfDependents ==> Findings and Treatmet

In [120]: